a **Accuracy against training cost**

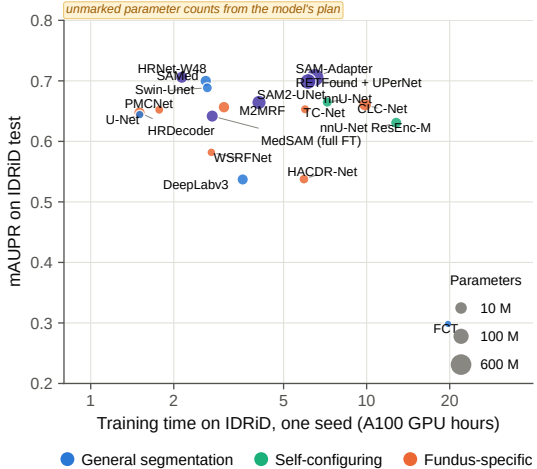

b **Resources per model**

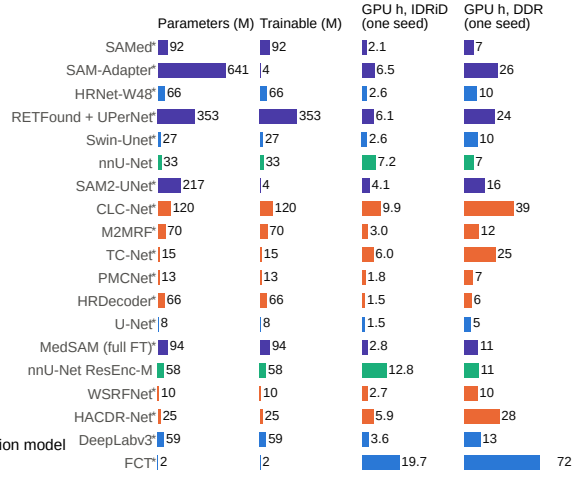

Models ordered by mAUPR on the IDRiD test set. 40 GB A100, one task per card; nnU-Net follows its own 1,000-epoch schedule. * parameter count read from the run (the others are the plan's); FLOPs, peak memory and inference time join when the cost audit has run.