a **Accuracy against training cost**

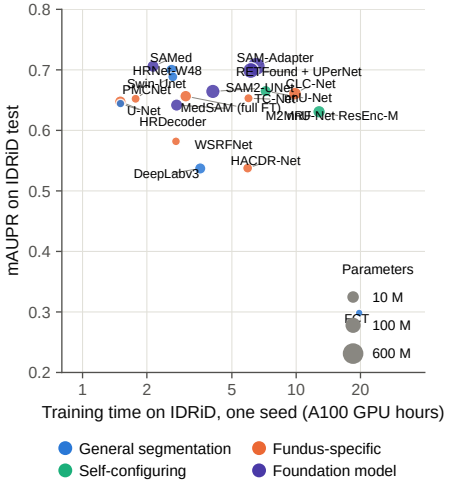

b **Resources per model**

	Parameters (M)	Trainable (M)	GPU h, IDRiD (one seed)	GPU h, DDR (one seed)	FLOPs (G), one canvas	Peak infer. memory (GB)	Inference (s/image)
SAMed*	92	4	2.1	7	–	1.9	0.30
SAM-Adapter*	641	4	6.5	26	11960	7.4	0.92
HRNet-W48*	66	66	2.6	10	984	1.3	0.05
RETFound + UPerNet*	353	353	6.1	24	20870	2.5	1.10
Swin-Unet*	27	27	2.6	10	1169	0.5	0.17
nnU-Net*	59	59	7.2	7	–	–	–
SAM2-UNet*	217	4	4.1	16	2677	1.6	0.21
CLC-Net*	120	120	9.9	39	11760	3.9	0.86
M2MRF*	70	70	3.0	12	931	1.1	0.05
TC-Net*	15	15	6.0	25	9631	8.7	0.29
PMCNet*	13	13	1.8	7	443	1.9	0.03
HRDecoder*	66	66	1.5	6	–	–	–
U-Net*	8	8	1.5	5	589	1.4	0.03
MedSAM (full FT)*	94	94	2.8	11	1951	5.5	0.19
nnU-Net ResEnc-M*	106	106	12.8	11	–	–	–
WSRFNet*	10	10	2.7	10	831	1.4	0.05
HACDR-Net*	25	25	5.9	28	–	–	–
DeepLabv3+*	59	59	3.6	13	2543	0.9	0.05
FCT*	2	2	19.7	72	2942	0.3	0.94

Models ordered by mAUPR on the IDRiD test set. 40 GB A100, one task per card; nnU-Net follows its own 1,000-epoch schedule.
* parameter count read from the run (the others are the plan's). FLOPs, peak inference memory and inference time: one 1,440-pixel canvas with the model's own windows, median of five passes on an A100; – not measured (HACDR-Net and HRDecoder: their mmcv forks do not build in the audit environment; nnU-Net: its own pipeline).