a **Accuracy against training cost**

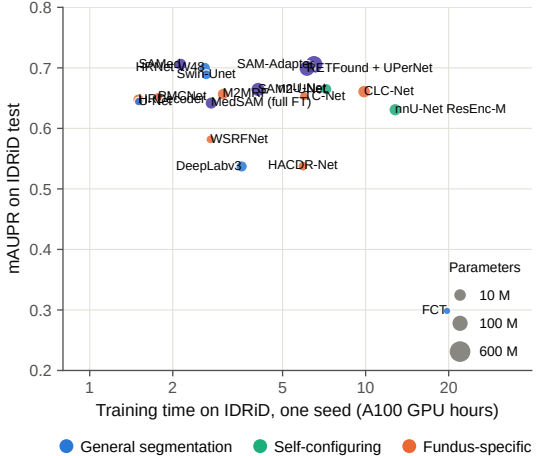

b **Resources per model**

	Parameters (M)	Trainable (M)	GPU h, IDRiD (one seed)	GPU h, DDR (one seed)
SAMed*	92	92	2.1	7
SAM-Adapter*	641	4	6.5	26
HRNet-W48*	66	66	2.6	10
RETFound + UPerNet*	353	353	6.1	24
Swin-Unet*	27	27	2.6	10
nnU-Net*	59	59	7.2	7
SAM2-UNet*	217	4	4.1	16
CLC-Net*	120	120	9.9	39
M2MRF*	70	70	3.0	12
TC-Net*	15	15	6.0	25
PMCNet*	13	13	1.8	7
HRDecoder*	66	66	1.5	6
U-Net*	8	8	1.5	5
MedSAM (full FT)*	94	94	2.8	11
nnU-Net ResEnc-M*	106	106	12.8	11
WSRFNet*	10	10	2.7	10
HACDR-Net*	25	25	5.9	28
DeepLabv3*	59	59	3.6	13
FCT*	2	2	19.7	72

Models ordered by mAUPR on the IDRiD test set. 40 GB A100, one task per card; nnU-Net follows its own 1,000-epoch schedule. * parameter count read from the run (the others are the plan's); FLOPs, peak memory and inference time join when the cost audit has run.