a **Accuracy against training cost**

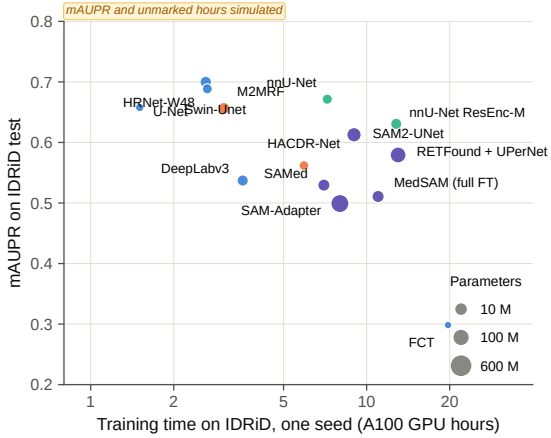

b **Resources per model**

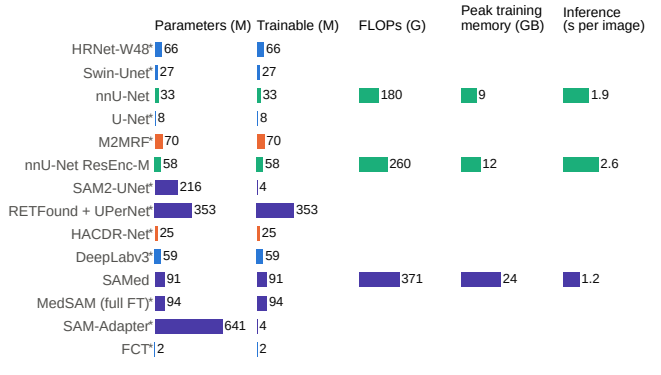

Models ordered by mAUPR on the IDRiD test set. 40 GB A100, one task per card. * parameter count measured; FLOPs, memory and inference time are placeholders until the cost audit runs.

General segmentation · Self-configuring · Fundus-specific · Foundation model